

Brief Announcement: PUSH, a DISC Shell

Noah Evans
Nara Institute of Science and Technology
Nara, Japan
noah-e@is.naist.jp

Eric Van Hensbergen
IBM Research
Austin, TX
bergevan@us.ibm.com

ABSTRACT

This paper explores the use of extended shell pipeline operators to establish distributed workflows and correlate results.

Categories and Subject Descriptors: C.2.4 [Network Operating Systems]: Inferno

General Terms: Design, Management

Keywords: Supercomputing, Shell, Distributed Pipeline

The deluge of huge data sets such those provided by sensor networks, online transactions, and the web provide exciting opportunities for data analytics. The scale of the data makes it impossible to deal with in a reasonable amount of time on isolated machines. This has lead to Data Intensive Supercomputing [2] emerging as the standard tool for solving research problems using these vast datasets. In typical DISC systems, runtimes [3] [1] [4] define graphs of processes, the edges of the graphs representing pipes and their vertices represent computation on a system. Within these runtimes a new class of languages [6] [7] [5] can be used by researchers to solve "pleasantly parallel" problems more quickly without worrying about explicit concurrency.

These languages provide automated control flow (typically matched to the architecture of the underlying runtime) and channels of communication between systems. In existing systems, these workflows and the underlying computation are tightly linked, tying solutions to a particular runtime, workflow and language. This creates difficulties for analytics researchers who wish to draw upon tools written in many different languages or runtimes which may be available on several different architectures or operating systems.

We observe that UNIX pipes were designed to get around many of these incompatibilities, allowing developers to hook together tools written in different languages and runtimes in ad-hoc fashions. This allowed tool developers to focus on doing one thing well, and enabled code portability and reuse in ways not originally conceived by the tool authors. The UNIX shell incorporated a model for tersely composing these smaller tools in pipelines (e.g. `'sort | uniq -c'`), creating coherent workflow to solve more complicated problems quickly. Tools read from standard input and wrote to standard output, allowing programs to work together in streams *with no explicit knowledge of this chaining built into the program itself*.

One to one pipelines such as those used by a typical UNIX

shell, can not be trivially mapped to DISC workflows which incorporate one-to-many, many-to-many, and many-to-one data flows. Additionally, typical UNIX pipeline tools write data according to buffer boundaries instead of record boundaries. As [6] notes DISC systems need to be able to cleanly separate input streams into records and then show that the order of these records is independent. By separating input and output into discrete unordered records data can be easily distributed and coalesced.

To address these issues we have implemented a prototype shell, which we call PUSH, using DISC principles and incorporating extended pipeline operators to establish distributed workflows and correlate results. Our implementation is based on extending an existing shell, MASH, from which we inherited a rich interpreted scripting language. It treats variables as lists of strings and has no native handling for any other data type. Integer expression handling and other facilities are provided by shell commands. It has native regular expression support and it has a novel ability to do declarative shell programming through a make like syntax incorporated in the shell itself.




Figure 1: The structure of the PUSH shell

We have added two additional pipeline operators, a multiplexing fan-out($|<[n]$), and a coalescing fan-in(>1). This combination allows PUSH to distribute I/O to and from multiple simultaneous threads of control. The fan-out argument n specifies the desired degree of parallel threading. If no argument is specified, the default of spawning a new thread per record (up to the limit of available cores) is used. This can also be overridden by command line options or environment variables. The pipeline operators provide implicit grouping semantics allowing natural nesting and composability. While their complimentary nature usually lead to symmetric mappings (where the number of fan-outs equal the number of fan-ins), there is nothing within our implementation which enforces it. Normal redirections as well as application specific sources and sinks can provide alternate

data paths. Remote thread distribution and interconnect are composed and managed using synthetic file systems in much the same manner as Xcpu, pushing the distributed complexity into the middleware in an language and runtime neutral fashion.

PUSH also differs from traditional shells by implementing native support for record based input handling over pipelines. This facility is similar to the argument field separators, IFS and OFS, in traditional shells which use a pattern to determine how to tokenize arguments. PUSH provides two variables, ORS and IRS, which point to record separator modules. These modules (called multiplexors in PUSH) split data on record boundaries, emitting individual records that the system distributes and coalesces.

The choice of which multipipe to target is left as a decision to the module. Different data formats may have different output requirements. Demultiplexing from a multipipe is performed by creating a many to one communications channel within the shell. The shell creates a reader processes which connects to each pipe in the multipipe. When the data reaches an appropriate record boundary a buffer is passed from the reader to the shell which then writes each record buffer to the output pipeline.

An example from our particular experience, Natural Language Processing, is to apply an analyzer to a large set of files, a "corpus". User programs go through each file which contain a list of sentences, one sentence per line. They then tokenize the sentence into words, finding the part of speech and morphology of the words that make up the sentence. This sort of task maps very well to the DISC model. There are a large number of discrete sets of data whose order is not necessarily important. We need to perform a computationally intensive task on each of the sentences, which are small, discrete records and ideal target for parallelization.

PUSH was designed to exploit this mapping. For example, to get a histogram of the distribution of Japanese words from a set of documents using chasen, a Japanese morphological analyzer, we take a set of files containing sentences and then distribute them to a cluster of machines on our network. The command is as follows:

```
push -c '{
  ORS=./blm.dis
  du -an files |< xargs os \\\
    chasen | awk '{print \$1}' | sort | uniq -c \\\
    >| sort -rn
}'
```

The first variable, ORS, declares our record multiplexor module, the intermediary used to ensure that the input and output to distributed pipes are correctly aligned to record boundaries. `du -n` gives a list of the files (note that our `du` is a bit different from the canonical UNIX `du`, it replaces much of find's functionality) which are then "fanned out" (`|<`) using a combination of a *multipipe*, an ordered set of pipes, and a *multiplexor* which determines which pipes are the targets of each unit of output. This fanned out data goes to `xargs` on other machines which then uses the filenames(sent from the instantiating machine) as arguments to `chasen`. The `du` acts as a command driver, fanning out file names to the individual worker machines. The workers then use the filenames input to `xargs`, which uses the input filenames as arguments to `xargs` target command. Using the output of the analyzer `awk` extracts the first line fields(Japanese words) which are

then sorted and counted using `uniq`. Finally these word counts are "fanned in" (`>|`) to the originating machine which then sorts them.

We currently have a working prototype of the PUSH shell, which we can use to target local distributed clusters, dynamic clusters built using Amazon's EC2 cloud, and large scale clusters such as a Blue Gene running the kittyhawk infrastructure. We are currently in the process of evaluating and optimizing performance for a variety of application types.

Because PUSH distributes streaming data over potentially remote links push is prone to causing network congestion between nodes. We are investigating multiple approaches to mitigating this overhead, both by using parallel distributed file systems and/or using distributed databases such as Google's Bigtable. One research challenge is to find a way to integrate reference to these external data resources (or record sets within these resources) with the distributed pipelines in a natural fashion.

Another area of ongoing exploration is looking at adding fault tolerance as well as a more automated mechanism for determining how many resources to allocate to a particular problem and automatically adapting that provisioning based on monitoring the changing workload as it executes on the cluster.

PUSH is under active development, but the current source code is open source and available on Google Code. This work has been supported by the Department of Energy Of Office of Science Operating and Runtime Systems for Extreme Scale Scientific Computation project under contract #DE-FG02-08ER25851.

1. REFERENCES

- [1] A. Bialecki, M. Cafarella, D. Cutting, and O. O Malley. Hadoop: a framework for running applications on large clusters built of commodity hardware, 2005. *Wiki at <http://lucene.apache.org/hadoop>*.
- [2] R.E. Bryant. Data-Intensive Supercomputing: The case for DISC. Technical report, Tech Report CMU-CS-07-128, Carnegie Mellon University, School of Computer Science, 2007.
- [3] J. Dean and S. Ghemawat. MapReduce: Simplified Data Processing on Large Clusters. *Communications of the ACM*, 51(01):7, 2008.
- [4] M. Isard, M. Budiu, Y. Yu, A. Birrell, and D. Fetterly. Dryad: distributed data-parallel programs from sequential building blocks. In *Proceedings of the 2007 conference on EuroSys*, pages 59–72. ACM Press New York, NY, USA, 2007.
- [5] C. Olston, B. Reed, U. Srivastava, R. Kumar, and A. Tomkins. Pig latin: a not-so-foreign language for data processing. In *Proceedings of the 2008 ACM SIGMOD international conference on Management of data*, pages 1099–1110. ACM New York, NY, USA, 2008.
- [6] R. Pike. Interpreting the data: Parallel analysis with Sawzall. *Scientific Programming*, 13(4):277–298, 2005.
- [7] Y. Yu, M. Isard, D. Fetterly, M. Budiu, U. Erlingsson, P.K. Gunda, and J. Currey. DryadLINQ: A system for general-purpose distributed data-parallel computing using a high-level language. In *Symposium on Operating System Design and Implementation (OSDI)*, San Diego, CA, December, pages 8–10, 2008.